

ANUJOT SINGH

+91 93129 41201 · anujott.singh@gmail.com · [linkedin.com/in/anujotsingh](https://www.linkedin.com/in/anujotsingh) · github.com/anujott-codes

EDUCATION

B.Tech — Artificial Intelligence & Machine Learning GGSIPU — USAR, New Delhi

CGPA: 9.1 / 10

2024-2028

TECHNICAL SKILLS

Languages: Python, C++, SQL, HTML/CSS

ML / DL: scikit-learn, XGBoost, LightGBM, PyTorch, TensorFlow, Transformers, DistilBERT, spaCy, Sentence-Transformers, SHAP, Optuna

Generative AI: LLMs, RAG Pipelines, LangChain, Prompt Engineering, Google Gemini API, Embeddings, Mistral AI, OpenAI, HuggingFace, Natural Language Processing (NLP), Semantic Searching, Cosine Similarity

Vector DBs: ChromaDB, FAISS

MLOps & Backend: MLflow, DVC, FastAPI, Pydantic, Docker, GitHub Actions CI/CD, GHCR

Deployment: Render, Railway, Google Cloud Platform (GCP)

PROJECTS

VidMind — AI Meeting Intelligence Platform | *FastAPI, LangChain, Mistral AI, Whisper, ChromaDB, HuggingFace Embeddings, Docker*

[GitHub ↗](#)

[Demo ↗](#)

- Developed an end-to-end AI-powered meeting intelligence platform that processes YouTube videos and uploaded recordings to generate structured summaries, action items, decisions, and intelligent conversational Q&A over meeting transcripts.
- Engineered a Retrieval-Augmented Generation (RAG) pipeline using LangChain, ChromaDB, sentence-transformer embeddings, and OpenAI Whisper for multilingual audio preprocessing, transcript generation, context-aware retrieval, and accurate response generation.
- Designed scalable NLP workflows using Mistral AI for long-form recording analysis, implementing Map-Reduce summarization and information extraction to efficiently process lengthy transcripts while maintaining response quality.
- Built a modular FastAPI-based system with reusable NLP, RAG, API, and processing services, optimized through token-aware chunking, vector embeddings, and MMR-based retrieval to improve relevance, deployment readiness, maintainability, and contextual accuracy.

ResumeRadar — Intelligent Resume-to-JD Matching Engine | *Python, FastAPI, Transformers, DistilBERT, spaCy, PyTorch, DVC, Docker*

[GitHub ↗](#)

[Demo ↗](#)

- Developed an AI-powered resume screening system that analyzes PDF resumes against job descriptions and generates ATS-style compatibility scores with detailed skill-gap analysis.
- Built an end-to-end NLP pipeline combining transformer-based Named Entity Recognition (NER), skill classification, semantic embeddings, and similarity scoring for contextual resume-job matching.
- Fine-tuned a DistilBERT classifier to categorize extracted skills as hard or soft skills, achieving **82.2% classification accuracy** on a held-out test dataset.
- Implemented semantic skill matching using BAAI BGE embeddings and cosine similarity, enabling detection of equivalent skills beyond exact keyword matches.
- Engineered a scalable FastAPI backend supporting PDF parsing, OCR fallback, resume segmentation, skill extraction, and structured JSON-based analysis results.
- Integrated DVC-based dataset versioning and Dockerized deployment workflows to improve experiment reproducibility and model management.

Churn-ML Platform — End-to-End Customer Churn MLOps Pipeline | *LightGBM, FastAPI, Optuna, MLflow, DVC, Docker, GitHub Actions*

[GitHub ↗](#)

[Demo ↗](#)

- Built a production-grade customer churn prediction platform featuring a six-stage ML pipeline (data ingestion, validation, feature engineering, hyperparameter tuning, model training, and evaluation), with business-driven churn features engineered from customer behavior signals including support interactions, payment delays, spending velocity, and engagement recency.
- Developed and optimized a LightGBM-based churn prediction model using Optuna-driven Bayesian hyperparameter tuning and stratified cross-validation to maximize ROC-AUC performance, alongside MLflow-based experiment tracking and model registry workflows enabling automated model versioning, artifact management, and champion-model deployment.
- Engineered a FastAPI inference service with real-time feature engineering, schema validation, batch prediction support, and MLflow model loading for production inference.
- Integrated DVC-based dataset versioning, Docker containerization, automated testing, and GitHub Actions CI/CD pipelines to establish a fully reproducible MLOps workflow.

Approv.io — Explainable Credit Risk Assessment | *XGBoost, SHAP, FastAPI, React, Scikit-learn*

[GitHub ↗](#)

- Developed an AI-powered credit and loan approval platform using XGBoost models served through FastAPI APIs for real-time eligibility prediction.
- Integrated SHAP explainability techniques to provide transparent feature-level explanations for automated approval decisions.
- Built a full-stack web application using React, Tailwind CSS, and FastAPI with interactive dashboards and real-time EMI calculations.
- Designed preprocessing and inference pipelines using Scikit-learn components for scalable and consistent model serving.

CERTIFICATIONS

Hugging Face Transformers Course · Docker & Azure Fundamentals — Microsoft Learn Student Ambassadors · **Hack2Skill × GDG**

ACHIEVEMENTS & ACTIVITIES

Smart India Hackathon (SIH) 2024 & 2025 — National-level participant · **Microsoft Learn Student Ambassador** — Conducted workshops on Docker & Cloud fundamentals at GGSIPU